

Semester	Course Code	Name of the course	L	T	P	Credits
V	CS5M004	Ethics, Bias and Safety in LLM	2	0	0	2

Course Outcomes:


Sr. No	Course Outcome number	CO statement
1	CO1	Students will be able to Analyze ethical challenges in LLM development and deployment.
2	CO2	Students will be able to Detect and mitigate bias in training datasets and model outputs.
3	CO3	Students will be able to Assess and improve the safety of LLM-based applications.
4	CO4	Students will be able to Understand legal and regulatory frameworks surrounding AI usage.
5	CO5	Students will be able to Evaluate and propose solutions to ethical controversies in LLMs using real-world case studies.

Unit II	Privacy, Misinformation, and Hallucinations: Data privacy concerns in LLM training and inference, Types and sources of misinformation and hallucinations, Mechanisms of hallucinations in generative models, Combatting fake news and deepfakes with AI.	05
---------	---	----

Data Privacy Concern in LLM Training and Inference



J D College of Engineering
& Management, Nagpur
An Autonomous Institute






Fig. 2. Security attack scenario (training/fine-tuning phase): The malicious user puts a prompt in the triggering word "CF", which activates the backdoor injected during the pre-training or fine-tuning phase. The LLM then generates the output desired by the malicious developer. Components: (1) attacker - malicious developer, (2) attack entity - training/fine-tuning data, model/algorithm, etc., and (3) attack goal - malfunctioning through generating harmful or inappropriate content.





Fig. 3. Privacy attack scenario (inference phase): The malicious user puts a jailbreaking prompt in a way, causing the LLM to generate the desired output. Components: (1) attacker - malicious LLM user, (2) attack entity - private training/fine-tuning data, malicious prompt & black-box access to the model, (3) attacker's goal - harmful/toxic/violent content generation, PII extracted from the model, sensitive information extraction from the model.



Gradient Leakage Attack

Deep learning models are often trained using optimization algorithms that involve gradients. Gradients represent the direction of the steepest increase in a function that helps to optimize model parameters during training to minimize the loss function. If attackers can access or infer these gradients or the gradient information, they may obtain access to the model or even compromise its privacy and safety, e.g., reconstructing private training data. Sensitive information can be extracted by analyzing gradients during the training or manipulating the training data

Data Privacy Concern in LLM Training and Inference



J D College of Engineering
& Management, Nagpur
An Autonomous Institute

Category	Name	Source Code	Characteristics	Limitation	Potential Mitigation
Gradient Leakage Attacks	TAG [55]	GitHub	- Performs gradient attack to recover the local training data on Transformer-based language model. - Impact: Moderate	- Adversary requires access to the gradient of the Transformer model.	- Noise Perturbation [331]. - Differential Privacy [2].
	LAMP [14]	GitHub	- Recovers training data from gradients in FL environment using guided search with auxiliary LM. - Impact: High	- Does not deal with reconstructing training labels.	- Differential Privacy [2]
	GDBA [81]	GitHub	- Introduced parameterized adversarial distribution to perform gradient-based text-attack. - The attack is transferable to other LMs and application for varieties of tasks. - Impact: Moderate	- The method is only able to recover tokens.	- Robust optimization in training so that adversarial perturbation has minimum impact [313].

Membership Inference Attack:

The primary goal of a Membership Inference Attack (MIA) is to determine if a data sample has been included in an ML model's training data. The attacker can execute MIAs even in the absence of direct access to the underlying ML model parameters, relying solely on the observation of its output. Typically, such attacks take advantage of models' tendency to overfit their training data, resulting in lower loss values for training samples

Data Privacy Concern in LLM Training and Inference



J D College of Engineering
& Management, Nagpur
An Autonomous Institute

Membership Inference Attacks	MIA on PLMs [294]	N/A	- Attacker collects a small number of training data of the model which the PLMs have pre-trained on and trains the attack model with these instances. - Designs a binary classifier (member or non-member), leads to MIA. - Impact: Moderate	- Attack model stands on an unrealistic assumption of having access to some sample of training data.	- Adversarial regularization [180]. - MemGuard [109].
	PREMIA [65]	N/A	- Performs MIA based on preferred and non-preferred responses against specific prompts in LLMs. - Based on the response of the model the attacker determines the membership of the training set. - Impact: Moderate	- Attack was fully evaluated on open-source LLMs. The performance might not be same on close-source model such as GPT-3.	- DP-SGD [2] - Model pruning[86]. - Knowledge distillation [92]
	SaMIA [119]	GitHub	- Introduces sampling-based pseudo likelihood (SPL) method for MIA based on the n-gram similarity between pre-fix (part of generated text) and reference text (remaining part of the generated text). - Impact: Moderate	- Attacks performed on LLMs which training data is known. - Performance might degrade on the unknown data on which model has been trained.	- To be studied.

PII Leakage Attack PII, refers to data that , either alone or in combination with other information ,can uniquely identify an individual .PII encompasses direct identifiers like passport details and quasi-identifiers such as race and date of birth . Sensitive PII includes information like name, phone number, address, social security number (SSN) , financial , and medical records , while non-sensitive PII ,is readily available in public sources, such as zip code, race, and gender. Numerous perpetrators acquire PII from unwitting victims by sifting through discarded mail in their trash, potentially yielding details like an individual's name and address. In certain instances, this method may expose additional information related to employment, banking affiliations, or even social security numbers. Phishing and social engineering attacks leverage deceitful websites or emails, employing tactics designed to deceive individuals into disclosing critical details such as names, bank account numbers, passwords, or SSNs. Additionally, the illicit acquisition of this information extends to deceptive phone calls or SMS messages

Data Privacy Concern in LLM Training and Inference



J D College of Engineering
& Management, Nagpur
An Autonomous Institute

PII Leakage Attacks	TAB [102]	N/A	- An evaluation technique if the PII can be leaked under any leaked under any threat model. - Impact: Low	- The proposed method was evaluated on RNN and Transformer-based language model, did not demonstrate applicability on LLM or other model architectures.	- DP-SGD [2] - API Hardening [102]
	PII Compass [178]	N/A	- Proposed PII leakage method by prepending hand-crafted template with true prefix subject to different data in the adversary set. - Impact: Low	- Limited to extracting only phone number due to the lack of publicly available PII entries such as SSN.	- Adversarial Training [178] - PII masking [157] - DP-SGD [2]
	ProPILE [122]	N/A	- Introduced as a tool for risk assessment of PII leakage from open pre-trained transformer language model. - Based on the linkability of the given information of asked query evaluates the risk of PII exposure. - Impact: Moderate	- The performance highly depends on the evaluation that is available on open-source datasets. - The evaluation dataset might contain noisy or misleading information.	- Cleaning PII with structured patterns with regular expression [122]
	Memorization [23]	GitHub	- Demonstrated the PII leakage attack due to memorization in LLMs, i.e., GPT-2. - The proposed attack efficiently extracts training samples using certain prefixes by black-box query access - Impact: Low	- The number of successful extraction of PII are extremely low. - The attack was evaluated on GPT-2, it may not perform well on the latest LLMs, e.g., Llama, GPT-4.	- Training Data curation[45]. - DP-SGD [2] - Prompt-tuning [188]
	KART [177]	GitHub	- Defines four factors that contributes to a privacy leakage in PLMs, i.e., attacker's prior knowledge about the model, the target information to be leaked, resources to attack, and the availability of the target information in pre-trained data. [122] - Impact: Low	- It did not provide any universal attack methods in which all the factors contributed to execute a successful attack. - The impact on attack performance after applying DP was not evaluated.	- Differential Privacy [2] - Limiting attacker access to training data.

An AI hallucination is when a large language model (LLM) powering an artificial intelligence (AI) system generates false information or misleading results, often leading to incorrect human decision-making. Hallucinations are most associated with LLMs, resulting in incorrect textual output. However, they can also appear in AI-generated video, images and audio. Another term for an AI hallucination is a *confabulation*.

There are many possible reasons for hallucinations in LLMs, including the following:

Poor data quality. Hallucinations might occur when there is bad, incorrect or incomplete information in the data used to train the LLM. LLMs rely on a large body of training data to be able to produce output that's as relevant and accurate to the user who provided the input prompt. However, this training data can contain noise, errors, biases or inconsistencies; consequently, the LLM produces incorrect and, in some cases, completely nonsensical outputs.

There are many possible reasons for hallucinations in LLMs, including the

Generation method. Hallucinations can also occur from the training and generation methods used, even when the data set itself is consistent and reliable, and contains high-quality training data. For example, the model's previous generations might create bias or the transformer might perform false decoding, both of which result in the system hallucinating its response. Models might also have a bias toward generic or specific words, which might influence the information they generate or lead them to fabricate their response.

Input context. If the input prompt provided by the human user is unclear, inconsistent or contradictory, hallucinations can arise. While the quality of training data and training methods used are out of the users' control, they can control the inputs they provide to the AI system. By honing their inputs and providing the right context to the AI system, they can get it to produce better results.

Types of AI hallucinations

There are several types of AI hallucinations, including the following:

Sentence contradiction. This is when an LLM generates a sentence that contradicts a previous sentence.

For example:

Prompt: Write a description of a landscape in four-word sentences.

Output: The grass was green. The mountains were blue. The river was purple. The grass was brown.

Prompt contradiction. This is when a sentence contradicts the prompt used to generate it.

For example:

Prompt: Write a birthday card for my niece.

Output: Happy anniversary, Mom and Dad!

Factual contradiction. This is when fictitious information is presented as a fact.

For example:

Prompt: Name three cities in the United States.

Output: New York, London, Toronto.

Irrelevant or random hallucinations. This is when random information with little or no relation to the input is generated.

Prompt: Describe London to me.

Output: London is a city in England. Cats need to be fed at least once a day.

Example of an AI hallucination

One infamous example of an AI hallucination event occurred in February 2023 when Google's chatbot, Gemini, made an incorrect claim about the James Webb Space Telescope (JWST). In a promotional video, Gemini was prompted, "What new discoveries from the James Webb Space Telescope can I tell my 9-year-old about?" Gemini responded that the JWST took the very first pictures of an exoplanet outside the Earth's solar system. That information was false, since the first images of an exoplanet were taken in 2004 by the European Southern Observatory's Very Large Telescope (VLT) and the JWST was not launched until 2021. Bard's answer sounded plausible and was consistent with the prompt provided, but was pointed out as false by news agency Reuters.

Example of an AI hallucination

Another AI hallucination occurred in late 2022 when Meta demoed Galactica, an open-source LLM trained on millions of pieces of scientific information from textbooks, papers, encyclopedias, websites and even lecture notes. Galactica was designed for use by science researchers and students, many of whom reported that the system generated inaccurate, suspicious or biased results. One professor at New York University said that Galactica failed to accurately summarize his work, while another user said that it invented citations and research when he prompted it to perform a literature review on the topic of whether HIV causes AIDS.

Types of generative models

Autoregressive models predict the next data point in a sequence based on previous data instances. Transformers excel at natural language processing (NLP) tasks due to their enhanced ability to process context.

Diffusion models create new data by gradually adding noise to a dataset, then figuring out how to remove the noise and yield novel output.

- **Natural language processing:** Transformers can process complex natural language queries and respond conversationally with automated text generation, making them ideal for use as chatbots. For example, ChatGPT is OpenAI's chatbot implementation of their GPT generative model. Other NLP applications include sentiment analysis, speech recognition, text-to-speech (TTS) applications and document summarization.
- **Coding support:** The same autoregressive capabilities that enable transformers to excel at text generation also allow them to debug code and generate code snippets.
- **Time-series forecasting:** Autoregression can be easily applied to time-series forecasting, in which a model predicts future values based on previous trends. Time-series forecasting is frequently applied to financial modeling, market predictions and weather forecasting.
- **Reinforcement learning:** Transformers have begun to see use in reinforcement learning, a machine learning training technique that teaches autonomous decision-making. Transformers are also being applied to classification tasks.

- **Image generation:** Diffusion models power mainstream image generation and image synthesis tools such as Midjourney, Stable Diffusion and OpenAI's DALL-E 2. These models generate images in response to user prompts. Diffusion models can generate high-quality realistic images, including those of human faces.
The US Copyright Office ruled in 2023 that AI-generated images are not entitled to copyright protections. Meanwhile, numerous ongoing lawsuits will eventually determine whether AI-generated images are considered copyright violations.
- **Inpainting and outpainting:** Inpainting is the process of adding or removing content within an image, while outpainting expands an image beyond its original borders.
- **3D modeling:** Google's DreamFusion and NVIDIA's Magic3D are diffusion models that create 3D models from text inputs.
- **Market research:** Diffusion models show how things evolve over time, making them useful for understanding how consumers react to a product.
- **Anomaly detection:** Because they can learn how data changes over time, diffusion models can identify when data points do not fit established trends. Anomaly detection applications include cybersecurity, fraud prevention and disease detection.

Types of generative models

Generative adversarial networks (GANs) pair a discriminative and generative model together in a competition, with the goal being for the generator to create output that fools the discriminator.

Variational autoencoders compress input data with an encoder, then reverse the process with a decoder to create new similar data

Flow-based models learn the relationships between simple and complex distributions of data through reversible mathematical operations

- **Computer vision:** Computer vision is the use of machine learning to process information from images. Common computer vision tasks include object detection and classification, facial recognition, sign language translation and object tracking.
- **Image generation:** GANs can outperform diffusion models in the generation of realistic images. They also require less training time and are more compute-efficient, though diffusion models offer finer control, greater versatility and more stability. The Diffusion-GAN framework trains a GAN with diffusion to maximize the benefits of both model types.
- **Anomaly detection:** When trained to generate normal datasets, GANs can be applied to anomaly detection tasks. The GAN creates a synthetic dataset modeled on the real-world data, and by comparing the two, anomalies in the latter stand out. The discriminator can also detect anomalies by declaring certain data points more likely to be fake.
- **Data augmentation:** Data augmentation—the use of pre existing data to create more data samples—can boost computer vision performance even further with CNNs. This process differs from synthetic data in that it expands on real data, as opposed to generating something from scratch.

- **Image generation:** VAEs see use in image generation, though with mainstream image generation applications, diffusion models have widely replaced them. Compared to other image generators, VAEs tend to produce blurrier images due to their “averaging” of the latent space.
- **Genomics:** VAEs aid geneticists in calculating breeding values—the projected value an animal will provide with its offspring—as well as assigning disease risk scores.
- **Anomaly detection:** VAEs are cheaper and easier to train than both GANs and diffusers, making them an attractive choice for anomaly detection tasks. The re-created data is compared to the original data to isolate instances that deviate from the projected distribution.
- **Data imputation:** VAEs can generate new data to replace missing data and restore corrupted files. Examples include clearing up audio files and denoising videos as well as medical imaging. While VAEs tend to generate blurry images from scratch, they can restore pre existing blurry images by denoising the image.
- **Semisupervised learning:** VAEs help train classifiers by capturing data distributions in datasets with incomplete labeling. VAEs can also perform data augmentation to generate extra training samples for the classifier.

- **Image generation:** Flow-based models generate images by running randomly sampled noise through normalization flows to produce a clean image. PixelCNN is a notable 2016 image generation model that applies autoregression to the sequence of transformations, modeling images pixel by pixel.
- **Accurate density estimation:** Density estimation is the challenge of identifying the probability density function of a dataset based on observable data. While flow-based models can accurately represent complex data distributions, they struggle to capture long-range dependencies when compared to transformers.
- **Molecular graph generation:** Flow-based models have found use in molecular graph generation, in which new molecular structures are modeled as graphs by generative AI models. Molecular graph generation is an important tool in aiding scientists and engineers in the discovery of new drugs and materials.

Inaccurate data sources.

Generative AI is highly reliant upon extensive data sets for training because it is primarily focused on skill set acquisition through guided pattern recognition. If the training data set provided is not accurate in any way, contains quantum biases, or is riddled with inconsistency

Prioritization of Fluency Over Accuracy

The majority of the existing generative AI models have been trained to produce outputs that are grammatically and fluently consistent rather than accurate. This manner of design may result in hallucinations when the model tries to systematically guess the missing information by filling in gaps with information that seemingly makes sense

Misinterpretation of Input

Generative AI can make mistakes with user queries, particularly if the user input is vague or overly complicated. In cases where the model does not understand the context in which the request is located nor the purpose of the request, it can provide a response that is loosely connected to the main issue yet misses the point entirely

Overgeneralization of Patterns

An AI model, in most cases, has a strength in detecting and copying patterns from the provided data. However, students could hallucinate by overgeneralizing from too few or anomalous examples.

Model Architecture

The design and architecture of the models themselves can contribute to hallucinations. For example, certain neural network architectures may be more prone to errors than others, especially when handling complex tasks

Lack of Common Sense

AI models lack true understanding or common-sense reasoning. They do not “know” facts in the way humans do. Therefore, they may generate information that lacks contextual coherence or logical consistency

Generative AI hallucinations can occur in multiple ways over different platforms. Below, we have discussed some of the examples.

- **Content Generation:** An AI tasked with writing a blog post about climate change might generate a section that gives fake reasons about some concept.
- **Image Generation:** A model meant to create more realistic-looking animals may produce images with anatomical inaccuracies, such as having an extra limb.
- **Fabricated Facts:** An AI might be asked to summarize a scientific study. Instead of accurately describing the findings, it could invent new details about a study that does not exist.
- **Business Analytics Blunders:** In a corporate setting, an AI system analyzing customer feedback might generate a report highlighting non-existent product features, leading to misguided strategic decisions.

These examples tell us how hallucination could differ from context and application.

Strategies to Manage Generative AI Hallucinations

Generative AI hallucinations elimination is best achieved through technical refinement and modification of user behavior practices.

1. Improving the Quality of Training Datasets

The main approach that reduces hallucinations is the use of precise, well-structured, and bias-free comprehensive datasets. Developers' strict adherence to these measures greatly enhances the chances of AI algorithms not learning fictitious narratives and content.

2. Domain-Specific Model Tuning

General models can be trained to use more accurate, specific, and comprehensible datasets aimed at specific tasks to reduce hallucinations. Specializing in medical or legal documents has proven to be helpful, as further restricting the model's focus increases accuracy within those fields.

3. Verification Methods Incorporation

Verification layers, such as cross-referencing the output with other sources or having another AI model double-check the information, can help remove hallucinatory content before the user sees it.

4. Providing Accurate Factual Prompts

The chance of hallucinations occurring can be minimized by ensuring that a user provides a model with constructive, clear, and concise instructions. Broad queries are much more likely to trigger hallucinations, while precise questions tend to yield more accurate outputs.

5. Verifying AI-Generated Content

The users must verify essential information sources generated by AI systems through established reliable external references. It is vital to verify facts since hallucination disinformation creates potential risks for making decisions based on incorrect or distorted information.

6. Developing Better Evaluation Metrics: The development of evaluation metrics represents a research demand because developers and researchers require tools for uncovering and evaluating hallucinations in AI-generated content. These evaluation methods serve to direct development work on models and their respective training and design methods.

7. Raising Awareness: Educating users about the limitations of generative AI and the possibility of hallucinations is crucial. Users should approach AI-generated content with a critical mindset, fact-checking information when necessary.

8. Continuous Research: Ongoing research in the field of AI is essential for developing better models and techniques to handle hallucinations. Innovations in algorithms and architectures can help minimize errors over time.

Misinformation and disinformation are commonly known as 'fake news'. The mis/disinformation phenomenon has been catalysed by the technological revolution. The increasingly frequent use of social media contributes to the amount of data shared and processed by algorithms. It has also changed the way news is created and disseminated.

Put simply, misinformation is the sharing of incorrect information whereas disinformation is the sharing of incorrect information with a **malicious intent**.

Mis/disinformation can take many forms such as memes, misattributed content, fabricated or cloned websites and news. It also includes altered and fabricated images and videos (these are also called deepfakes).

Bots, trolls and other AI or non-AI agents can increase the reach of mis/disinformation campaigns. Algorithms that create filter bubbles significantly intensify its impact and effect on users.

Fake news: “Fictions deliberately fabricated and presented as non-fiction with the intent to mislead recipients into treating fiction as fact or into doubting verifiable fact.” Eric Emanuelson, Jr., Fake Left, Fake Right: Promoting an Informed Public in the Era of Alternative Facts, 70 Admin. L. Rev. 209, 218-19 (2018).

Deepfakes: Media created with machine learning and/or artificial intelligence in which a person in an existing image or video is replaced with someone else’s likeness. Viorela Dan et al., Visual Mis- and Disinformation, Social Media, and Democracy, 98 Journalism & Mass Communication Quarterly 641, 643 (2021).

Misinformation: A “verifiably false statement of fact.” Charles R. Corbett, Chemtrails and Solar Geoengineers: Governing Online Conspiracy Theory Misinformation, 85 Mo. L. Rev. 633, 652 (2020).

Disinformation: “[F]alse information that is purposely spread to deceive people.” David M. J. Lazer et al., The science of fake news, 359 Science 1094 (2018).

Information literacy: A set of “competencies that enable people to critically and effectively engage with communications content; the institutions that facilitate this content; and the use of digital technologies.” Media and information literate citizens: think critically, click wisely!, UNESCO Digital Library.

Psychological Operation (PSYOP): “Often confused with Information Operations as a whole, PSYOP refers to influence activities specifically intended 'to induce or reinforce foreign attitudes and behavior in a manner favorable to U.S. objectives.'” 12 (March 8, 2011) Terrorist Use of the Internet: Information Operations in Cyberspace

How Do Deepfakes Work?

Deepfake applications work in various ways. Some transfer the facial movements of an actor to a target video, such as the one we saw at the beginning of this article, or this Obama deepfake created by comedian Jordan Peele to warn about the threat of fake news:

Watch this video : <https://youtu.be/cQ54GDm1eL0>

Other deepfakes map the face of a target person onto other videos—for example, this video of Nicolas Cage's face mapped onto that of characters in different movies.

Watch this video : <https://youtu.be/BU9YAHigNx8>

Deepfakes uses a special type of neural-network structure called an “autoencoder.” Autoencoders are composed of two parts: an encoder, which compresses an image into a small amount of data; and a decoder, which decompresses the compressed data back into the original image. The mechanism is similar to those of image and video codecs such as JPEG and MPEG

Deepfake applications use two autoencoders—one trained on the face of the actor and the other trained on the face of the target.

The application swaps the inputs and outputs of the two autoencoders to transfer the facial movements of the actor to the target.




Image courtesy of Two Minute Papers

First, you have to gather many photos of the faces of the target and the actor, and those photos must show each face from different angles. The process usually involves grabbing thousands of frames from videos that feature the target and actor and cropping them to contain only the faces. New deepfake tools such as Faceswap can do part of the legwork by automating the frame extraction and cropping, but they still require manual tweaking.

Deepfakes and the Technology Behind It from BBC

<https://youtu.be/8Lhl-e2B8Lg>

1. Don't jump to conclusions

One general best practice is to “refrain from jumping to strong conclusions in either direction without additional information and confirmation,” said Hany Farid, a professor of computer science at Dartmouth College and one of the world’s leading image forensics experts. Wait and see what other information comes to light about video or piece of content in question, and hold back on sharing it with others. This is especially true in breaking news situations, where information moves quickly and is often wrong or misinterpreted in the early hours.

2. Consider the source

Pay attention to where this piece of content originated. Was it uploaded to a random account on a social network? Who is claiming ownership of it, and where do they say it came from?

3. Check where else it is (and isn't) online

Do an online search with the video’s title or other relevant keywords to see if multiple credible news outlets and the person or entity in question have weighed in on it. Again, no need to rush to a conclusion until you can consult multiple courses.

4. Inspect the mouth

Farid said that today's deepfake videos often have telltale signs. One thing to look closely at is the mouth of the person speaking, because these tools often struggle "to accurately render the teeth, tongue, and mouth interior. So, keep a watchful eye on the mouth for [visual anomalies]."

5. Slow it down

Farid also encourages people to slow down and freeze parts of the video to watch it more closely. "Slowing a video down so that you can clearly see the frame-to-frame transition can sometimes reveal temporal glitches that are introduced in manipulated video," he said.

Watch this video :

https://video-player.buzzfeed.com/embed?video_id=54375

Types and Sources of misinformation and Hallucination



J D College of Engineering
& Management, Nagpur
An Autonomous Institute

